

Hand gesture recognition model that can accurately identify and classify different hand gestures from image or video data, enabling intuitive human-computer interaction and gesture-based control systems.

Importing libraries

```
In [4]: import numpy as np
import pandas as pd
import os
import cv2

import tensorflow as tf
from sklearn.model_selection import train_test_split
from keras.preprocessing.image import ImageDataGenerator

from tensorflow import keras
from keras.models import Sequential
from keras.layers import Dense, Dropout, Flatten
from keras.layers import Conv2D, MaxPooling2D

import matplotlib.pyplot as plt
```

Reading images

```
In [5]: dir = r'C:\Users\Satoshi\Desktop\Data\prodigy-4\leapGestRecog'
```

```
In [6]: images = []
labels = []
for directory in os.listdir(dir):
    for subDir in os.listdir(os.path.join(dir,directory)):
        for img in os.listdir(os.path.join(dir, directory, subDir)):
            img_path = os.path.join(dir, directory, subDir, img)
            images.append(img_path)
            labels.append(subDir)
```

```
In [7]: # images = np.array(images)
# labels = np.array(labels)
# labels
```

Converting the data to DataFrame

```
In [8]: Iseries = pd.Series(images, name="Images")
Lseries = pd.Series(labels, name="labels")
hand_gesture_data = pd.concat([Iseries, Lseries], axis=1)
hand_gesture_df = pd.DataFrame(hand_gesture_data)
```

counting the images in each class

```
In [10]: pd.Series(labels).value_counts()
```

```
Out[10]: 05_thumb      2000
01_palm      2000
03_fist      2000
07_ok        2000
02_l         2000
09_c         2000
08_palm_moved 2000
10_down      2000
04_fist_moved 2000
06_index     2000
dtype: int64
```

Splitting the dataset into train and test

```
In [11]: X_train, X_test = train_test_split(hand_gesture_df, test_size=0.2, random_state=42)
train_set, val_set = train_test_split(hand_gesture_df, test_size=0.3, random_state=42)
```

Data Preprocessing

```
In [12]: image_gen = ImageDataGenerator(preprocessing_function= tf.keras.applications.mobilenet_v2.preprocess_input)
train = image_gen.flow_from_dataframe(dataframe= train_set,x_col="Images",y_col="labels",
                                     target_size=(244,244),
                                     color_mode='rgb',
                                     class_mode="categorical",
                                     batch_size=4,
                                     shuffle=False
                                     )
test = image_gen.flow_from_dataframe(dataframe= X_test,x_col="Images", y_col="labels",
                                    target_size=(244,244),
                                    color_mode='rgb',
                                    class_mode="categorical",
                                    batch_size=4,
                                    shuffle= False
                                    )
val = image_gen.flow_from_dataframe(dataframe= val_set,x_col="Images", y_col="labels",
                                   target_size=(244,244),
                                   color_mode= 'rgb',
                                   class_mode="categorical",
                                   batch_size=4,
                                   shuffle=False
                                   )
```

Found 14000 validated image filenames belonging to 10 classes.

Found 4000 validated image filenames belonging to 10 classes.

Found 6000 validated image filenames belonging to 10 classes.

```
In [13]: classes=list(train.class_indices.keys())
print (classes)
```

```
['01_palm', '02_l', '03_fist', '04_fist_moved', '05_thumb', '06_index', '07_ok', '08_palm_moved', '09_c', '10_do
wn']
```

Data visualization


```
In [14]: def show_hand_gesture(image_gen):
test_dict = test.class_indices
classes = list(test_dict.keys())
images, labels=next(image_gen)
plt.figure(figsize=(20,20))
length = len(labels)
if length<25:
    r=length
else:
    r=25
for i in range(r):
    plt.subplot(5,5,i+1)
    image=(images[i]+1)/2
    plt.imshow(image)
    index=np.argmax(labels[i])
    class_name=classes[index]
    plt.title(class_name, color="green",fontsize=16)
    plt.axis('off')
plt.show()
```

```
In [15]: show_hand_gesture(train)
```


08_palm_moved




09_c



08_palm_moved



04_fist_moved



Building my model

```
In [20]: model = keras.models.Sequential([
keras.layers.Conv2D(filters=128, kernel_size=(8, 8), strides=(3, 3), activation='relu', input_shape=(224, 224, 3))
```

```

keras.layers.BatchNormalization(),

keras.layers.Conv2D(filters=256, kernel_size=(5, 5), strides=(1, 1), activation='relu', padding="same"),
keras.layers.BatchNormalization(),
keras.layers.MaxPool2D(pool_size=(3, 3)),

keras.layers.Conv2D(filters=256, kernel_size=(3, 3), strides=(1, 1), activation='relu', padding="same"),
keras.layers.BatchNormalization(),
keras.layers.Conv2D(filters=256, kernel_size=(1, 1), strides=(1, 1), activation='relu', padding="same"),
keras.layers.BatchNormalization(),
keras.layers.Conv2D(filters=256, kernel_size=(1, 1), strides=(1, 1), activation='relu', padding="same"),
keras.layers.BatchNormalization(),

keras.layers.Conv2D(filters=512, kernel_size=(3, 3), activation='relu', padding="same"),
keras.layers.BatchNormalization(),
keras.layers.MaxPool2D(pool_size=(2, 2)),

keras.layers.Conv2D(filters=512, kernel_size=(3, 3), activation='relu', padding="same"),
keras.layers.BatchNormalization(),

keras.layers.Conv2D(filters=512, kernel_size=(3, 3), activation='relu', padding="same"),
keras.layers.BatchNormalization(),

keras.layers.MaxPool2D(pool_size=(2, 2)),

keras.layers.Conv2D(filters=512, kernel_size=(3, 3), activation='relu', padding="same"),
keras.layers.BatchNormalization(),

keras.layers.MaxPool2D(pool_size=(2, 2)),

keras.layers.Flatten(),
keras.layers.Dense(1024, activation='relu'),
keras.layers.Dropout(0.5),
keras.layers.Dense(1024, activation='relu'),
keras.layers.Dropout(0.5),
keras.layers.Dense(10, activation='softmax')
])

model.compile(
    loss='categorical_crossentropy',
    optimizer=tf.optimizers.SGD(learning_rate=0.001),
    metrics=['accuracy']
)

model.summary()

```

Model: "sequential_2"

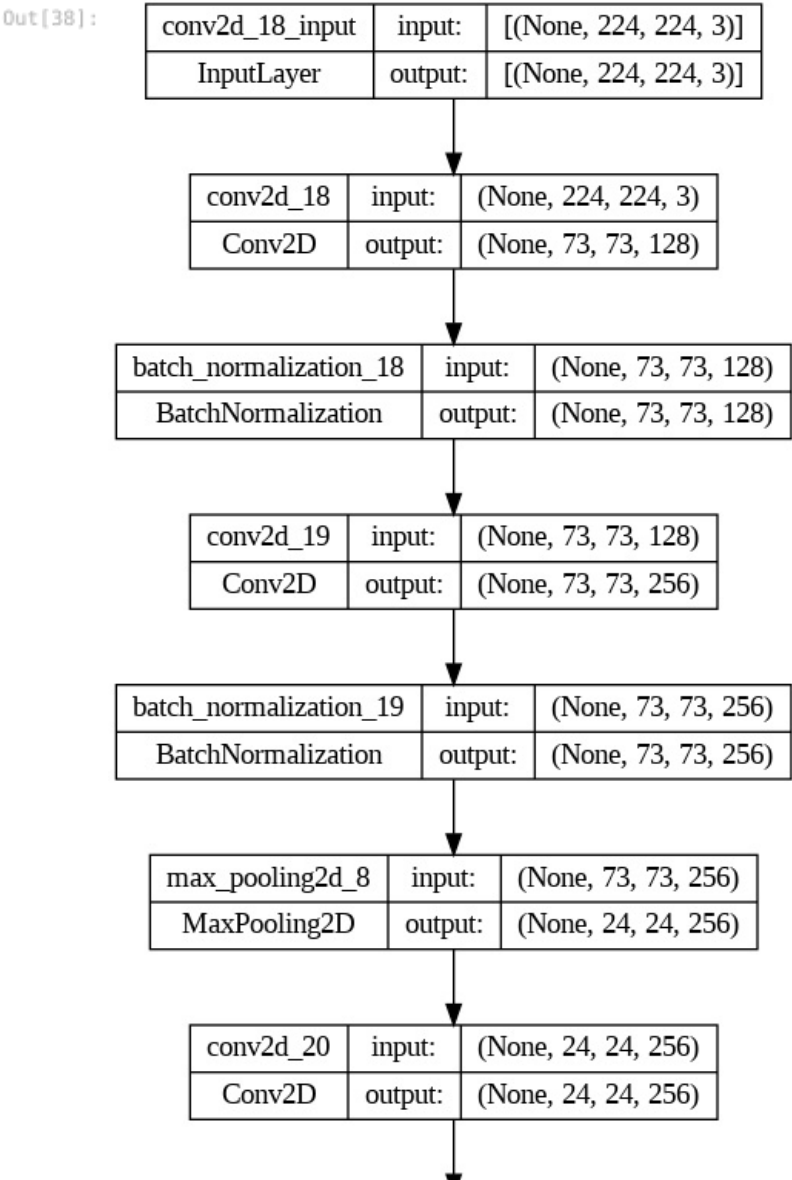
Layer (type)	Output Shape	Param #
=====		
conv2d_18 (Conv2D)	(None, 73, 73, 128)	24704
batch_normalization_18 (Batch Normalization)	(None, 73, 73, 128)	512
conv2d_19 (Conv2D)	(None, 73, 73, 256)	819456
batch_normalization_19 (Batch Normalization)	(None, 73, 73, 256)	1024
max_pooling2d_8 (MaxPooling2D)	(None, 24, 24, 256)	0
conv2d_20 (Conv2D)	(None, 24, 24, 256)	590080

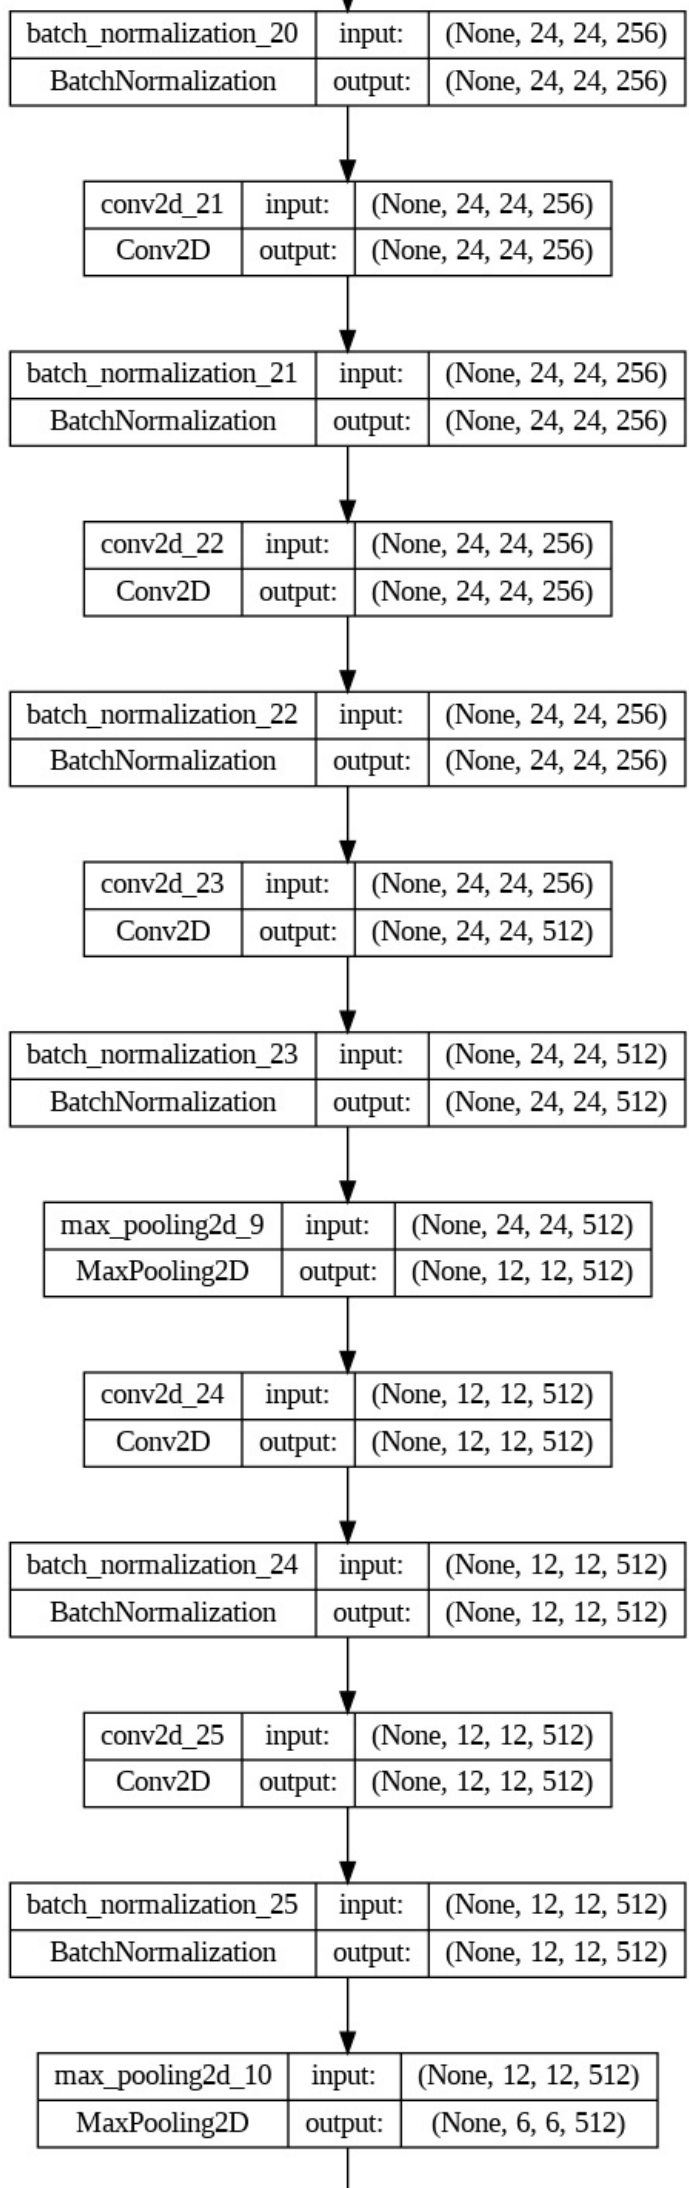
batch_normalization_20 (Batch Normalization)	(None, 24, 24, 256)	1024
conv2d_21 (Conv2D)	(None, 24, 24, 256)	65792
batch_normalization_21 (Batch Normalization)	(None, 24, 24, 256)	1024
conv2d_22 (Conv2D)	(None, 24, 24, 256)	65792
batch_normalization_22 (Batch Normalization)	(None, 24, 24, 256)	1024
conv2d_23 (Conv2D)	(None, 24, 24, 512)	1180160
batch_normalization_23 (Batch Normalization)	(None, 24, 24, 512)	2048
max_pooling2d_9 (MaxPooling2D)	(None, 12, 12, 512)	0
conv2d_24 (Conv2D)	(None, 12, 12, 512)	2359808
batch_normalization_24 (Batch Normalization)	(None, 12, 12, 512)	2048
conv2d_25 (Conv2D)	(None, 12, 12, 512)	2359808
batch_normalization_25 (Batch Normalization)	(None, 12, 12, 512)	2048
max_pooling2d_10 (MaxPooling2D)	(None, 6, 6, 512)	0
conv2d_26 (Conv2D)	(None, 6, 6, 512)	2359808
batch_normalization_26 (Batch Normalization)	(None, 6, 6, 512)	2048
max_pooling2d_11 (MaxPooling2D)	(None, 3, 3, 512)	0

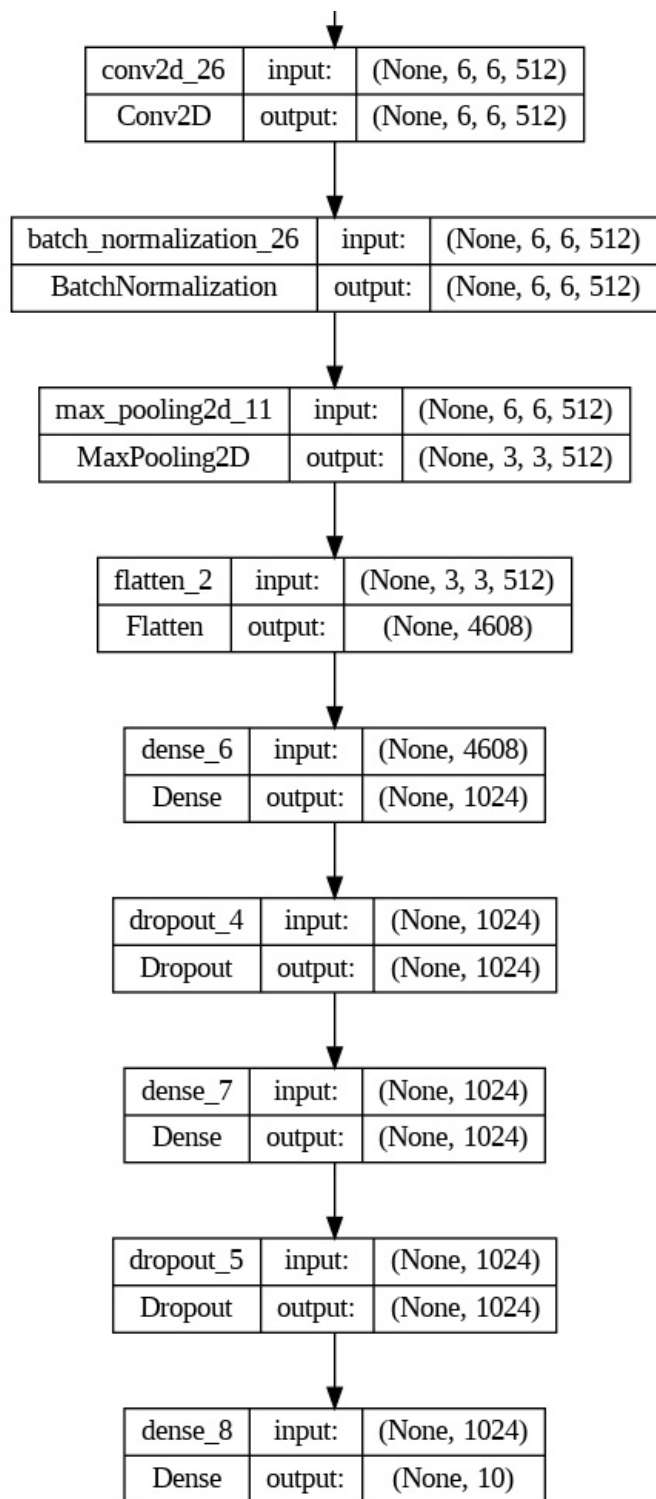
flatten_2 (Flatten)	(None, 4608)	0
dense_6 (Dense)	(None, 1024)	4719616
dropout_4 (Dropout)	(None, 1024)	0
dense_7 (Dense)	(None, 1024)	1049600
dropout_5 (Dropout)	(None, 1024)	0
dense_8 (Dense)	(None, 10)	10250

```
=====
Total params: 15617674 (59.58 MB)
Trainable params: 15611274 (59.55 MB)
Non-trainable params: 6400 (25.00 KB)
```

```
In [38]: from keras.utils import plot_model
plot_model(model, to_file='model_plot.png', show_shapes=True, show_layer_names=True)
```







Training my mdoel

```
In [21]: history = model.fit(train, epochs=3, validation_data=val, verbose=1)
```

Epoch 1/3

3500/3500 [=====] - 197s 56ms/step - loss: 0.4896 - accuracy: 0.8483 - val_loss: 0.0043 - val_accuracy: 0.9993

Epoch 2/3

3500/3500 [=====] - 194s 56ms/step - loss: 0.0190 - accuracy: 0.9944 - val_loss: 0.0011 - val_accuracy: 0.9997

Epoch 3/3

3500/3500 [=====] - 204s 58ms/step - loss: 0.0077 - accuracy: 0.9981 - val_loss: 0.0012 - val_accuracy: 0.9995

Testing my model

```
In [22]: model.evaluate(test, verbose=1)
```

```
1000/1000 [=====] - 18s 18ms/step - loss: 0.0013 - accuracy: 0.9995
```

```
Out[22]: [0.0013108871644362807, 0.9994999766349792]
```

Saving my model

```
In [23]: model.save("hand_gesture_Model.h5")
```

```
/usr/local/lib/python3.10/dist-packages/keras/src/engine/training.py:3103: UserWarning: You are saving your model as an HDF5 file via `model.save()`. This file format is considered legacy. We recommend using instead the native Keras format, e.g. `model.save('my_model.keras')`.
```

```
saving_api.save_model(
```


Getting results

```
In [24]: pred = model.predict(test)
pred = np.argmax(pred, axis=1)


labels = (train.class_indices)
labels = dict((v,k) for k,v in labels.items())
pred2 = [labels[k] for k in pred]
```

```
1000/1000 [=====] - 20s 20ms/step
```

```
In [25]: plt.plot(history.history['accuracy'])
plt.plot(history.history['val_accuracy'])
plt.title('model accuracy')
plt.ylabel('accuracy')
plt.xlabel('epoch')
plt.legend(['train', 'val'], loc='upper left')
plt.show()
```



```
In [26]: plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])
plt.title('model loss')
plt.ylabel('loss')
plt.xlabel('epoch')
plt.legend(['train', 'val'], loc='upper left')
plt.show()
```

```
In [32]: from sklearn.metrics import confusion_matrix, accuracy_score, classification_report
```

```
y_test = X_test.labels
print(classification_report(y_test, pred2))
print("Accuracy of the Model:", "{:.1f}%".format(accuracy_score(y_test, pred2)*100))
```

	precision	recall	f1-score	support
01_palm	1.00	1.00	1.00	404
02_l	1.00	1.00	1.00	377
03_fist	1.00	1.00	1.00	404
04_fist_moved	1.00	1.00	1.00	410
05_thumb	1.00	1.00	1.00	417
06_index	1.00	1.00	1.00	366
07_ok	1.00	1.00	1.00	418
08_palm_moved	1.00	1.00	1.00	403
09_c	1.00	1.00	1.00	392
10_down	1.00	1.00	1.00	409
accuracy			1.00	4000
macro avg	1.00	1.00	1.00	4000
weighted avg	1.00	1.00	1.00	4000

Accuracy of the Model: 100.0%

```
In [37]: import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.metrics import confusion_matrix

class_labels = ['Palm', 'l', 'Fist', 'Fist_moved', 'Thumb', 'Index', 'Ok', 'Palm_moved', 'C', 'Down']

cm = confusion_matrix(y_test, pred2)

plt.figure(figsize=(10, 5))
sns.heatmap(cm, annot=True, fmt='g', vmin=0, cmap='Blues')

plt.xticks(ticks=[0.5, 1.5, 2.5, 3.5, 4.5, 5.5, 6.5, 7.5, 8.5, 9.5], labels=class_labels, rotation=25)
plt.yticks(ticks=[0.5, 1.5, 2.5, 3.5, 4.5, 5.5, 6.5, 7.5, 8.5, 9.5], labels=class_labels, rotation=0)
plt.xlabel("Predicted")
```

```
plt.ylabel("Actual")  
plt.title("Confusion Matrix")  
plt.show()
```

